



Welcome to the New York Restaurant Guide,
An analysis of Food Safety inspections and machine learning model

Today's Fare:

Department of Health Inspection Data

*\$
Gratis*

- Public Data provided by New York Open Data
- 3 year period from mid 2016 to mid 2019
- Scored and Graded inspections for food safety
 - A < 13 points

The New York City Department of Health and Mental Hygiene make regular inspections of all restaurants in New York City and approximately 3 years of results are made publicly available through the New York Open Data Project. Food Safety Inspections are scored based on the number and type of health code violations and grade is awarded. A score of 13 or less is awarded an A grade.




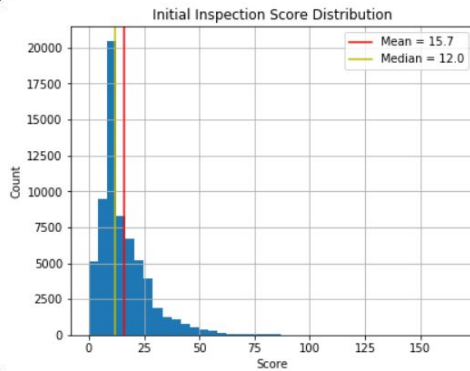
The meal this evening consists of an analysis of inspections results, scores and violations

This is followed by a deconstruction of the restaurant characteristics into 3 categories, it's location, type and history and their relationship with score and grade

Next are predictive grade classifier models available in 2 constructions

We finish with conclusions and business recommendations and a discussion of future work

First Course

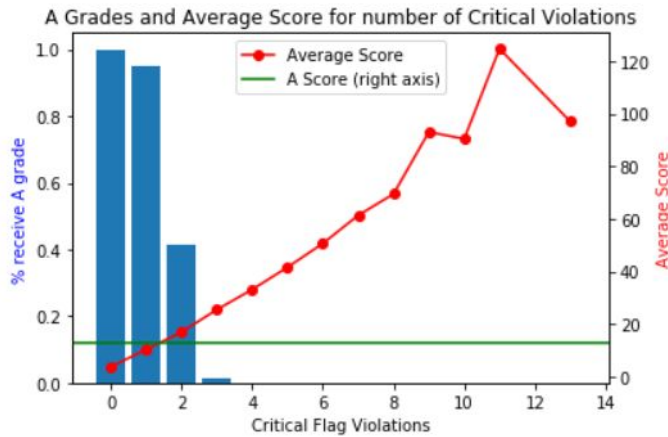


60% of initial inspections receive an A grade

Initial inspection scores follow an approximately normal distribution with a left skew and long tail that drags up the mean.

Approximately 60% of all initial inspections receive an A grade

First Course



1 Critical flag violation “allowable” for A

Critical flag violations present a serious risk to public health but a single violation will not get you shut down

95% of inspections that were cited for 1 critical violation were awarded an A grade

2 violations this number drops to 40%

Any more and it is virtually impossible to achieve an A

First Course


Correlation > 0.3

Critical Flags and Score

08A - Vermin Proofing

04L - Mice

02G - Cold food > 41 deg F



The heaviest correlation with inspection score is with the number of critical violations as these are codependent variables

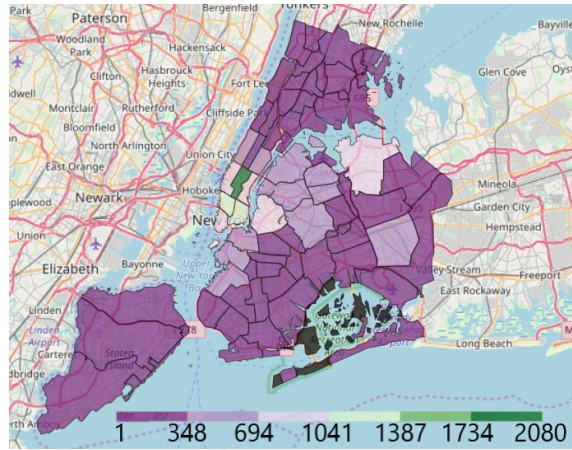
Next is a lack of sufficient vermin proofing in the establishment which is NOT a critical violation

Evidence of mice is the most common form of vermin cited by inspectors

Inadequate refrigeration of cold food items can cause serious illness and is scored highly for individual violations

Second Course


Community Boards
Highest Density in Midtown and
Lower Manhattan



New York City is divided up into 69 Community Board Areas
The highest density of restaurants is found in Midtown and Lower Manhattan

Second Course

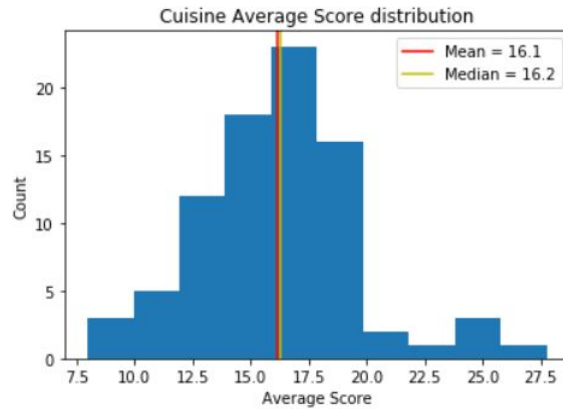
Community Boards
Average Inspection Score By
Community Board



The average inspection score by community board can be used to paint a picture of food safety standards across New York

Second Course

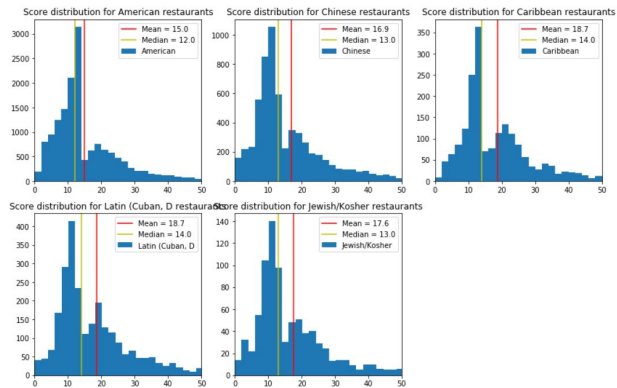
Cuisine Style
Average Inspection Score
Distribution by Cuisine Style



The relationship between cuisine style and score also follows an approximately normal distribution with certain cuisine styles much more easily able to consistently achieve a lower inspection score and others struggling

Second Course

Cuisine Style
Bimodal Score Distribution
Within Popular Cuisine Styles




The most popular cuisine styles with many inspection data points show a bimodal distribution.

Inspectors use their judgement or owners apply pressure and there is some hesitancy to give score that just misses an A

Second Course

Chain
Restaurants Belonging to Chains
Have a Lower Average Score

Score distribution for chain and not chain restaurants WITHOUT outliers

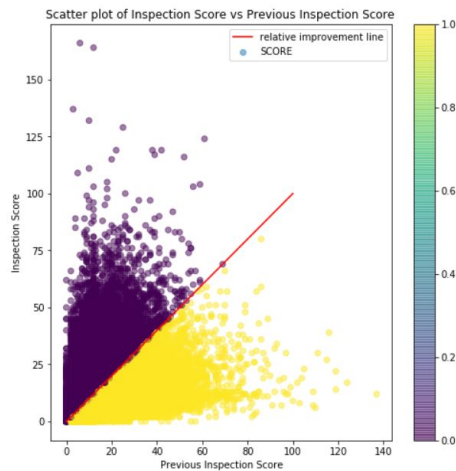


Chains with more than one outlet operating under the same business name are most able to score lower on their inspection. Larger chains often have strict training programs and have very clear processes based around food safety.

Second Course

History
Improvement and
disimprovement equally likely

Dramatically bad scores do not
happen again

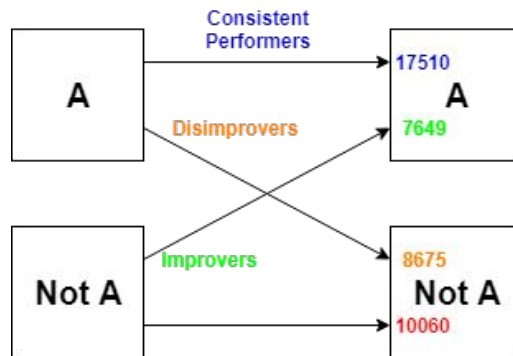


Restaurants both improve and disimprove over the course of multiple inspections but dramatically bad scores are addressed seriously by the Dept of Health and process changes are enforced. They do not happen again

Second Course

History
A Rated Restaurants Most
Commonly Maintain Their A

Improvement and
Disimprovement Possible and
Commonplace




An A is the most common inspection outcome and restaurants that achieve an A will usually maintain that A through the next inspection cycle but a significant portion do not. It is not possible to rest on the laurels of a successful A. Likewise, a significant portion of restaurants that do not achieve an A use that as an impetus to improve and achieve an A on their next inspection.

Third Course

Categorical Matrix
Cuisine and Community Board
One-Hot Encoded

156 Features

	precision	recall	f1-score	support
False	0.69	0.66	0.67	5110
True	0.68	0.71	0.69	5149
accuracy			0.68	10259
macro avg	0.68	0.68	0.68	10259
weighted avg	0.68	0.68	0.68	10259



Using features from the 3 characteristics above 2 machine learning classifiers were created

The first uses one hot encoding of categorical features for community board and cuisine style to create a large sparse matrix with 156 columns consisting mostly of zeros

An accuracy rate of 68% is achieved at predicting an A grade

Third Course


Summary Statistics Matrix

Cuisine and Community Board
Features Using Mean, Standard
Deviation and Count

9 Features

Much Faster Training

	precision	recall	f1-score	support
False	0.68	0.64	0.66	6368
True	0.67	0.70	0.68	6456
accuracy			0.67	12824
macro avg	0.67	0.67	0.67	12824
weighted avg	0.67	0.67	0.67	12824



The second classifier uses the mean, standard deviation and count for each of the categories within cuisine style and community board to create numerical features and is a much smaller matrix. This classifier achieved an accuracy of 67% in one tenth of the training time for the optimised classifier



Advance knowledge of the expected inspection grade allows restaurateurs to identify their risk of not receiving an A and restaurant consultants to identify those same establishments where their services will be the most useful. 88% of New Yorkers are influenced by posted inspection grades when choosing where to dine. Getting an A matters.



The classifier models can be improved by more data. Inspections happen every single day and the Open Data database updates daily

Time of year effects are not considered in this analysis or this model but may be observed

Inspection history is considered only to a depth of one inspection, many restaurants have been operating for several years and have a much deeper inspection history that is not taken into account. Similarly new restaurants that have not had a previous health inspection since opening are not able to be considered in the classifier model.



Thank you for dining with us this evening!